

```
import os

import json

import time

from collections import Counter

from typing import List, Tuple


import joblib

import numpy as np

import pandas as pd

import streamlit as st

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.metrics import (

    confusion_matrix,

    roc_curve,

    auc,

    precision_recall_curve,

    PrecisionRecallDisplay,

)


st.set_page_config(page_title="Spam Classifier — Visualizations", layout="wide")
```

```
def ts() -> str:
```

```
    return time.strftime("%Y%m%d-%H%M%S")
```

```
@st.cache_data(show_spinner=False)
```

```
def load_csv(path: str) -> pd.DataFrame:
```

```
    return pd.read_csv(path)
```

```
@st.cache_data(show_spinner=False)
```

```
def list_datasets() -> List[str]:
```

```
    paths: List[str] = []
```

```
    for root in ("datasets", os.path.join("datasets", "processed")):
```

```
        if os.path.isdir(root):
```

```
            for name in os.listdir(root):
```

```
                p = os.path.join(root, name)
```

```
                if name.lower().endswith(".csv") and os.path.isfile(p):
```

```
                    paths.append(p)
```

```
    return sorted(paths)
```

```
def infer_cols(df: pd.DataFrame) -> Tuple[str, str]:
```

```
    # Try to infer label/text columns
```

```

cols = list(df.columns)

label_candidates = [c for c in cols if c.lower() in ("label", "target", "col_0")]

text_candidates = [c for c in cols if c.lower() in ("text", "message", "text_clean",
"col_1")]

label = label_candidates[0] if label_candidates else cols[0]

text = text_candidates[0] if text_candidates else cols[-1]

return label, text

```

```

def token_topn(series: pd.Series, topn: int) -> List[Tuple[str, int]]:

    counter: Counter = Counter()

    for s in series.astype(str):

        counter.update(s.split())

    return counter.most_common(topn)

```

```

@st.cache_resource(show_spinner=False)

def load_artifacts(models_dir: str):

    vec = joblib.load(os.path.join(models_dir, "spam_tfidf_vectorizer.joblib"))

    clf = joblib.load(os.path.join(models_dir, "spam_logreg_model.joblib"))

    pos, neg = "spam", "ham"

    meta_p = os.path.join(models_dir, "spam_label_mapping.json")

    if os.path.exists(meta_p):

        try:

```

```

        with open(meta_p, "r", encoding="utf-8") as f:

            meta = json.load(f)

            pos = meta.get("positive", pos)

            neg = meta.get("negative", neg)

    except Exception:

        pass

    return vec, clf, pos, neg

```

```

def label_to_int(series: pd.Series, pos_label: str = "spam") -> np.ndarray:

    s = series.astype(str).str.lower()

    return (s == pos_label.lower()).astype(int).values

```

Lightweight normalization to match training text_clean behavior for live inference

```

import re

URL_RE = re.compile(r"https?://\S+|www\.\S+", re.IGNORECASE)

EMAIL_RE = re.compile(r"\b[\w\.-]+@[\w\.-]+\.[a-zA-Z]{2,}\b")

PHONE_RE = re.compile(r"\b(?:\+?\d[\d\-\s]{7,}\d)\b")

```

```

def normalize_text(text: str, keep_numbers: bool = False) -> str:

    if not isinstance(text, str):

        text = "" if text is None else str(text)

    t = text.lower()

```

```

t = URL_RE.sub("<URL>", t)

t = EMAIL_RE.sub("<EMAIL>", t)

t = PHONE_RE.sub("<PHONE>", t)

if not keep_numbers:

    t = re.sub(r"\d+", "<NUM>", t)

t = re.sub(r"^[^\s<>]", " ", t)

t = re.sub(r"\s+", " ", t).strip()

return t


def main():

    st.title("Spam/Ham Classifier — Phase 4 Visualizations")

    st.caption("Interactive dashboard for data distribution, token patterns, and
model performance")


    # Sidebar: data and artifacts

    with st.sidebar:

        st.header("Inputs")

        datasets = list_datasets()

        ds_path = st.selectbox("Dataset CSV", datasets,
index=datasets.index("datasets/processed/sms_spam_clean.csv") if
"datasets/processed/sms_spam_clean.csv" in datasets else 0)

        df = load_csv(ds_path)

        label_col, text_col = infer_cols(df)

        label_col = st.selectbox("Label column", options=list(df.columns),
index=list(df.columns).index(label_col))

```

```
text_col = st.selectbox("Text column", options=list(df.columns),
index=list(df.columns).index(text_col))
```

```
models_dir = st.text_input("Models dir", value="models")
```

```
test_size = st.slider("Test size", min_value=0.1, max_value=0.4, value=0.2,
step=0.05)
```

```
seed = st.number_input("Seed", min_value=0, value=42, step=1)
```

```
threshold = st.slider("Decision threshold", min_value=0.1, max_value=0.9,
value=0.5, step=0.01)
```

```
st.subheader("Data Overview")
```

```
c1, c2 = st.columns(2)
```

```
with c1:
```

```
st.write("Class distribution")
```

```
counts = df[label_col].value_counts().sort_index()
```

```
st.bar_chart(counts)
```

```
with c2:
```

```
st.write("Token replacements in cleaned text (approximate)")
```

```
sample = df[text_col].astype(str)
```

```
repl = {
```

```
    "<URL>": sample.str.count(r"<URL>").sum(),
```

```
    "<EMAIL>": sample.str.count(r"<EMAIL>").sum(),
```

```
    "<PHONE>": sample.str.count(r"<PHONE>").sum(),
```

```
    "<NUM>": sample.str.count(r"<NUM>").sum(),
```

```
}
```

```

st.table(pd.DataFrame.from_dict(repl, orient="index", columns=["count"]))

st.subheader("Top Tokens by Class")

topn = st.slider("Top-N tokens", min_value=10, max_value=40, value=20,
step=5)

col_a, col_b = st.columns(2)

for label, col in [(counts.index[0], col_a), (counts.index[-1], col_b)]:

    with col:

        st.write(f"Class: {label}")

        top = token_topn(df.loc[df[label_col] == label, text_col], topn)

        if top:

            toks, freqs = zip(*top)

            fig, ax = plt.subplots(figsize=(6, 4))

            sns.barplot(x=list(freqs), y=list(toks), ax=ax, palette="viridis")

            ax.set_xlabel("frequency"); ax.set_ylabel("token")

            st.pyplot(fig)

        else:

            st.info("No tokens found.")

# Model-based visuals

st.subheader("Model Performance (Test)")

if os.path.exists(os.path.join(models_dir, "spam_tfidf_vectorizer.joblib")) and
os.path.exists(os.path.join(models_dir, "spam_logreg_model.joblib")):

    vec, clf, pos_label, neg_label = load_artifacts(models_dir)

    X = df[text_col].astype(str).fillna("")

```

```

y = label_to_int(df[label_col], pos_label=pos_label)

Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=test_size,
random_state=seed, stratify=y)

Xte_vec = vec.transform(Xte)

proba = clf.predict_proba(Xte_vec)[: , 1]

pred = (proba >= threshold).astype(int)

# Confusion matrix

cm = confusion_matrix(yte, pred)

cm_df = pd.DataFrame(cm, index=["true_0", "true_1"],
columns=["pred_0", "pred_1"])

st.write("Confusion matrix")

st.dataframe(cm_df)

# ROC/PR curves

fpr, tpr, _ = roc_curve(yte, proba)

roc_auc = auc(fpr, tpr)

prec, rec, _ = precision_recall_curve(yte, proba)

pr_fig, pr_ax = plt.subplots(1, 2, figsize=(10, 4))

pr_ax[0].plot(fpr, tpr, label=f"AUC={roc_auc:.3f}")

pr_ax[0].plot([0,1],[0,1], linestyle="--", color="gray")

pr_ax[0].set_title("ROC")

pr_ax[0].set_xlabel("FPR"); pr_ax[0].set_ylabel("TPR")

PrecisionRecallDisplay(precision=prec, recall=rec).plot(ax=pr_ax[1])

pr_ax[1].set_title("Precision-Recall")

```



```
st.pyplot(pr_fig)
```

```
# Threshold sweep small table
```

```
st.write("Threshold sweep (precision/recall/f1)")
```

```
ths = np.round(np.linspace(0.3, 0.8, 11), 3)
```

```
rows = []
```

```
for t in ths:
```

```
    p = (proba >= t).astype(int)
```

```
    from sklearn.metrics import precision_score, recall_score, f1_score
```

```
    rows.append({
```

```
        "threshold": t,
```

```
        "precision": float(precision_score(yte, p, zero_division=0)),
```

```
        "recall": float(recall_score(yte, p, zero_division=0)),
```

```
        "f1": float(f1_score(yte, p, zero_division=0)),
```

```
    })
```

```
st.dataframe(pd.DataFrame(rows))
```

```
# Live Inference
```

```
st.subheader("Live Inference")
```

```
# Provide two quick examples to try
```

```
ex_spam = "Free entry in 2 a wkly comp to win cash now! Call +44 906-170-1461 to claim prize"
```

```
ex_ham = "Ok, I'll see you at 7 pm for dinner. Thanks!"
```

```

c_ex1, c_ex2 = st.columns(2)

with c_ex1:

    if st.button("Use spam example"):

        st.session_state["input_text"] = ex_spam

with c_ex2:

    if st.button("Use ham example"):

        st.session_state["input_text"] = ex_ham


# Text area bound to session_state so examples populate it

if "input_text" not in st.session_state:

    st.session_state["input_text"] = ""

user_text = st.text_area("Enter a message to classify", key="input_text")


if st.button("Predict"):

    if user_text.strip():

        cleaned = normalize_text(user_text)

        with st.expander("Show normalized text", expanded=False):

            st.code(cleaned)

        X_single = vec.transform([cleaned])

        prob = float(clf.predict_proba(X_single)[:, 1][0])

        pred_label = pos_label if prob >= threshold else neg_label

        st.success(f"Prediction: {pred_label} | spam-prob = {prob:.4f}"
(threshold = {threshold:.2f}))

```

```

        # Probability bar (0..1) with threshold marker

        fig_g, ax_g = plt.subplots(figsize=(6, 0.6))

        ax_g.barh([0], [prob], color="#d62728" if pred_label == pos_label
else "#1f77b4")

        ax_g.axvline(threshold, color="black", linestyle="--", linewidth=1)

        ax_g.set_xlim(0, 1)

        ax_g.set_yticks([])

        ax_g.set_xlabel("spam probability")

        ax_g.text(min(prob + 0.02, 0.98), 0, f"{prob:.2f}", va="center")

        st.pyplot(fig_g)

    else:

        st.info("Please enter a non-empty message.")

```

```

else:

    st.info("Model artifacts not found in 'models/'. Train the model first to
enable performance plots.")

```

```

if __name__ == "__main__":

    main()

```